

✓ LSTM for Text Generation

Long Short-Term Memory (LSTM) is a special type of Recurrent Neural Network designed to overcome the limitations of vanilla RNNs. LSTM uses memory cells and gating mechanisms (input, forget, and output gates) to selectively store and update information over long sequences. In text generation, an LSTM learns long-term dependencies between characters or words and predicts the next token based on both recent and distant context, resulting in more coherent and less repetitive text compared to basic RNNs.

```
# =====
# LSTM BASED TEXT GENERATION
# =====

import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Embedding

# -----
# DATASET (SAME AS BEFORE)
# -----
text = """
artificial intelligence is transforming modern society.
it is used in healthcare finance education and transportation.
machine learning allows systems to improve automatically with experience.
data plays a critical role in training intelligent systems.
large datasets help models learn complex patterns.
deep learning uses multi layer neural networks.
neural networks are inspired by biological neurons.
each neuron processes input and produces an output.
training a neural network requires optimization techniques.
gradient descent minimizes the loss function.

natural language processing helps computers understand human language.
text generation is a key task in nlp.
language models predict the next word or character.
recurrent neural networks handle sequential data.
lstm and gru models address long term dependency problems.
however rnn based models are slow for long sequences.

transformer models changed the field of nlp.
they rely on self attention mechanisms.
attention allows the model to focus on relevant context.
transformers process data in parallel.
this makes training faster and more efficient.
modern language models are based on transformers.

education is being improved using artificial intelligence.
intelligent tutoring systems personalize learning.
automated grading saves time for teachers.
online education platforms use recommendation systems.
technology enhances the quality of learning experiences.

ethical considerations are important in artificial intelligence.
fairness transparency and accountability must be ensured.
ai systems should be designed responsibly.
data privacy and security are major concerns.
researchers continue to improve ai safety.

text generation models can create stories poems and articles.
they are used in chatbots virtual assistants and content creation.
generated text should be meaningful and coherent.
evaluation of text generation is challenging.
human judgement is often required.

continuous learning is essential in the field of ai.
research and innovation drive technological progress.
students should build strong foundations in mathematics.
programming skills are important for ai engineers.
practical experimentation enhances understanding.
"""

text = text.lower().replace("\n", " ")

# -----
# CHARACTER TOKENIZATION
# -----
chars = sorted(list(set(text)))
char_to_index = {c: i for i, c in enumerate(chars)}
index_to_char = {i: c for c, i in char_to_index.items()}
```

```

vocab_size = len(chars)
seq_length = 40

# -----
# CREATE INPUT-OUTPUT SEQUENCES
# -----
X, y = [], []

for i in range(len(text) - seq_length):
    X.append([char_to_index[c] for c in text[i:i+seq_length]])
    y.append(char_to_index[text[i+seq_length]])

X = np.array(X)
y = np.array(y)

# -----
# BUILD LSTM MODEL
# -----
model = Sequential([
    Embedding(vocab_size, 64),
    LSTM(128),
    Dense(vocab_size, activation='softmax')
])

model.compile(
    loss='sparse_categorical_crossentropy',
    optimizer='adam'
)

# -----
# TRAIN MODEL
# -----
model.fit(X, y, epochs=15, batch_size=64)

# -----
# TEXT GENERATION FUNCTION
# -----
def generate_text(seed_text, length=300):
    output = seed_text
    for _ in range(length):
        seq = np.array([[char_to_index[c] for c in output[-seq_length:]]])
        prediction = model.predict(seq, verbose=0)
        next_char = index_to_char[np.argmax(prediction)]
        output += next_char
    return output

# -----
# GENERATE TEXT
# -----
seed = "artificial intelligence "
generated_text = generate_text(seed)

print("Generated Text using LSTM:\n")
print(generated_text)

```

```

Epoch 1/15
34/34 ━━━━━━━━━━ 3s 9ms/step - loss: 3.1986
Epoch 2/15
34/34 ━━━━━━ 1s 7ms/step - loss: 2.8842
Epoch 3/15
34/34 ━━━━ 0s 6ms/step - loss: 2.8762
Epoch 4/15
34/34 ━━━━ 0s 6ms/step - loss: 2.7801
Epoch 5/15
34/34 ━━━━ 0s 5ms/step - loss: 2.6873
Epoch 6/15
34/34 ━━━━ 0s 5ms/step - loss: 2.5745
Epoch 7/15
34/34 ━━━━ 0s 13ms/step - loss: 2.5452
Epoch 8/15
34/34 ━━━━ 0s 6ms/step - loss: 2.4758
Epoch 9/15
34/34 ━━━━ 0s 7ms/step - loss: 2.3771
Epoch 10/15
34/34 ━━━━ 0s 5ms/step - loss: 2.3308
Epoch 11/15
34/34 ━━━━ 0s 5ms/step - loss: 2.2822
Epoch 12/15
34/34 ━━━━ 0s 6ms/step - loss: 2.2434
Epoch 13/15
34/34 ━━━━ 0s 6ms/step - loss: 2.1982
Epoch 14/15

```

```
34/34 ━━━━━━━━ 0s 5ms/step - loss: 2.1278
```

```
Epoch 15/15
```

```
34/34 ━━━━━━━━ 0s 5ms/step - loss: 2.0714
```

```
Generated Text using LSTM:
```

```
artificial intelligence and conters the monters the monters the monters the monters the monters the monters the
```

Limitations of LSTM

- LSTM models are computationally expensive compared to simple RNNs.
- Training LSTMs requires more time and memory due to complex gate operations.
- Performance depends heavily on the size and quality of the training data.
- LSTMs process sequences sequentially, which limits parallelization.
- For very long sequences, transformers outperform LSTM-based models.